

# Natural Language Document Classification of Philosophy Texts based on High-Level Qualitative Labels

Trent Yarosevich<sup>1</sup>

## Abstract

This paper uses document classification with a sample of academic philosophy papers with the aim of exploring whether these techniques are effective at identifying documents based on extremely high-level, qualitative labels that cover extremely large divisions within the philosophy community. In particular, the presumably easier task of distinguishing Continental and Analytic philosophy papers, which draw from very distinct lexicons, is considered first. Following this, we then try to distinguish 'good' from 'bad' philosophy papers, using samples drawn from some of the top Philosophy journals according to google scholar's citation analytics, which are labeled as 'good', and comparing them with papers drawn from graduate and undergraduate journals, which are labeled as 'bad'.  
[https://github.com/tyarosevich/AMATH\\_582/blob/master/Project/yarosevich582\\_project\\_final.pdf](https://github.com/tyarosevich/AMATH_582/blob/master/Project/yarosevich582_project_final.pdf)

<sup>1</sup>Department of Applied Mathematics, University of Washington, Seattle

## Contents

<b>1</b>	<b>Introduction and Overview</b>	<b>1</b>
<b>2</b>	<b>Theoretical Background</b>	<b>2</b>
2.1	Bag-of-words and Text Vectorization . . . . .	2
2.2	Classifiers - K-Nearest Neighbors and Discriminant Analysis . . . . .	2
	K-Nearest Neighbors • LDA and QDA	
<b>3</b>	<b>Algorithm Implementation and Development</b>	<b>3</b>
3.1	Data Selection . . . . .	3
3.2	Data Pre-processing . . . . .	3
3.3	Classification . . . . .	4
3.4	Visualization . . . . .	4
<b>4</b>	<b>Computational Results</b>	<b>4</b>
<b>5</b>	<b>Summary and Conclusions</b>	<b>4</b>
	<b>Appendix 1 - Python Functions</b>	<b>6</b>
	<b>Appendix 2 - Code Listing</b>	<b>6</b>

## 1. Introduction and Overview

Natural language processing (NLP) and machine-learning (ML) have become prevalent topics as algorithms and processing power have become adequate to the extremely complex task of analyzing written documents with thousands of different words. The simple size of the vocabulary is just a small part of the information contained in language, however; it also encodes concepts that are so far-reaching and abstract that even humans can struggle to rigorously define them. Philosophy has, among other things, historically concerned itself with rigorously defining concepts - from Plato's allegory of

the cave, to Hegel's introduction to *The Science of Logic*, to Gilles Deleuze's claim that philosophy *is* the study of the concept of concepts<sup>1</sup>. This is by no means an overture to a philosophical investigation, but merely serves as an observation that Philosophy uses language in a very abstracted way that makes use of rich syntactical information.

Classification of documents with NLP, on the other hand, often just considers individual words based on their frequency. This is called the 'bag-of-words' model, and this is the approach I will use here. The question is, can NLP with the bag-of-words approach classify texts whose labels are determined by extremely complex, abstract philosophical categories? To determine this, I will evaluate journal papers in two parts, one that will attempt to distinguish papers from two extremely broad fields within philosophy, and one that will attempt to distinguish between impactful papers that are widely read within the field, and papers that are written by students.

Before moving on to discuss the data acquisition and algorithms, a brief note is in order regarding the two broad fields in question: 'continental' and 'analytic' philosophy. Volumes have been written on what separates these two fields, but it makes for an interesting machine-learning task since they both have exclusive figures, as well as ones that overlap; exclusive lexicons of words, as well as common argumentative language; and exclusive topics as well as differing approaches to the same topics. I, of course, will not engage in any such comparison, but rather simply observe that even within Philosophy, the so-called 'continental-analytic divide' is both widely recognized and debated, and thus the question of whether or not they can be classified with machine-learning is at least

<sup>1</sup>Deleuze, Gilles, et al. What Is Philosophy? Verso, 2015.

interesting, regardless of whether it is a success or failure.

## 2. Theoretical Background

### 2.1 Bag-of-words and Text Vectorization

A fair amount of pre-processing was required for this investigation, but the most important overall aspect is the conversion of a text into numbers that faithfully represent that text so that they can be processed by a computer. There are many ways to do this, but the approach used here is called the 'bag-of-words' method, so-called since the arrangement of the individual words is discarded, and their frequency within each document is the only thing that matters. Thus in the preprocessing stage, the text must be reduced to a list of each word, and then a dictionary of all the unique words that occur in all the documents is formed. Each document is then represented by a word-count. For example, suppose we have three documents: "if, then", "if, or", "if, else". The resulting bag of words would look like this:

	"if"	"then"	"else"	"or"
Document 1	1	1	0	0
Document 2	1	0	0	1
Document 3	1	0	1	0

**Table 1.** An example bag-of-words matrix

The bag-of-words approach has a few issues, however, that need to be addressed by data-preprocessing, or 'text-cleaning'. One is the presence of 'noise' in the data in the form of what are called stop-words, or words that are so common and necessary for everyday language that they carry no information about the classification of the document. These words, such as 'the' and 'and', can simply be removed from the corpus. Another problem is term-frequency (TF). If word counts alone were used, the very common terms would overshadow the interesting ones that actually characterize the class to which the document belongs, that is to say their term-frequency would be very high. If a table like Table 1 is indexed as a matrix, then each term-frequency would be:

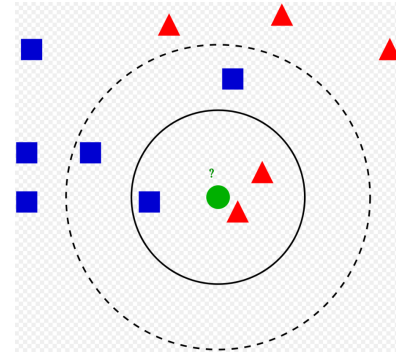
$$tf_{i,j} = \frac{n_{i,j}}{\sum_k n_{i,j}} \quad (1)$$

To avoid this problem, we use a method known as "inverse-document frequency," which is a weighting adjustment applied to the term-frequencies to give the TF-IDF vectors. This weighting is calculated as follows, where  $n$  is the number of documents and  $df(t)$  is the number of documents containing the particular term<sup>2</sup>:

$$idf(t) = \log \frac{1+n}{1+df(t)} + 1 \quad (2)$$

These TF-IDF values are then normalized with a Euclidean norm, and the resulting data matrix is then the set of features that will be passed to a classifier.

<sup>2</sup>[https://scikit-learn.org/stable/modules/feature\\_extraction.html#text-feature-extraction](https://scikit-learn.org/stable/modules/feature_extraction.html#text-feature-extraction)



**Figure 1.** KNN example in which the green dot is new data, and the squares and triangles are two classes in the training set. The inner circle represents  $n = 3$ , and the outer  $n = 5$ .

### 2.2 Classifiers - K-Nearest Neighbors and Discriminant Analysis

The results of pre-processing will be classified with three methods, K-Nearest Neighbors (KNN), Linear/Quadratic Discriminant Analysis (LDA/QDA), and Support Vector Classification (SVC). The latter will be treated as a black-box and will not be discussed here as it is simply presented as another point of comparison for accuracy. KNN and LDA/QDA will be briefly discussed below.

#### 2.2.1 K-Nearest Neighbors

The KNN approach is a very simple, naive supervised machine-learning method that requires labeled data. This algorithm's approach is simple: a training set of labeled data is used to evaluate any test data by finding the distance between each point in the new data-set and the labelled points in the training data-set. The distance  $d$  is defined as follows, in which  $\Delta x_n$  represents the difference in each direction between the test point and a particular neighbor:

$$d = \sqrt{\Delta x_1 + \Delta x_2 + \dots + \Delta x_n} \quad (3)$$

The test data point is then given the same label as its nearest training data point. This concept can also be extended by increasing the number of neighbors used to label the test data-point, and thus an odd number of neighbors (to avoid ties) can be evaluated to determine the label of the new point. This means that, if  $n = 3$ , the nearest neighbor might actually be ignored if the next two nearest neighbors belong a different class. Figure 1 provides a simple illustration of the algorithm<sup>3</sup>.

#### 2.2.2 LDA and QDA

The other classifying tool I'll discuss is LDA, which largely hinges upon the singular value decomposition, or SVD. The SVD is discussed widely in many textbooks and online resources, so for the sake of brevity I will not discuss it here, and instead focus on the approach taken by discriminant analysis simply noting that the decomposition of a data matrix  $A$

<sup>3</sup>[https://en.wikipedia.org/wiki/K-nearest\\_neighbors\\_algorithm#/media/File:KnnClassification.svg](https://en.wikipedia.org/wiki/K-nearest_neighbors_algorithm#/media/File:KnnClassification.svg)

has the following notation:

$$A = U\Sigma V^* \quad (4)$$

. With this approach, the orthogonal bases of the SVD are studied in order to try to understand what the data does when projected onto these bases, which eliminates their redundancy (co-variance). In essence, the LDA proceeds by examining the principal orthogonal modes, or POD, which are obtained from the SVD. We can then study how strong these modes are in the data set by looking at a particular sample's projection onto  $\Sigma V^*$ . The hope is then that different datapoints will have common features in a particular mode that can then be used to distinguish them from one another and thereby classify them.

Once a suitable mode is found that distinguishes labeled classes of data, linear discriminant analysis is used to project the data onto a subspace that both (a) maximizes the distance between the labeled classes of data, (b) minimizes the intra-class distance between points within a single labeled class. This is an optimization problem in which the between and within class matrices are given by:

$$S_B = (\mu_2 - \mu_1)(\mu_2 - \mu_1)^T$$

$$S_W = \sum_{j=1}^2 \sum_x (x - \mu_j)(x - \mu_j)^T \quad (5)$$

These values are then used for the optimization problem:

$$w = \operatorname{argmax}_w \frac{w^T S_B w}{w^T S_W w} \quad (6)$$

Finally, this value is used to solve the eigenvalue problem  $S_B w = \lambda S_W w$  in which  $\lambda$  gives the mode of interest and the associated eigenvector is the projection basis. In the subsequent analysis, this method will be handled with built-in methods of the sklearn library.

### 3. Algorithm Implementation and Development

#### 3.1 Data Selection

The data was downloaded manually, but in a more scaled-up version of this experiment, a web-scraper would be used, and so I will include the details here. The most important aspect of this is choosing a criteria that corresponds to the data labels, which can be difficult when the labels represent such high-level conceptual information - and by 'high-level' I mean the labels represent categories of philosophy that encompass a large number of other categories, and so on.

The principal tool used to select relevant papers published by active academics was the Google Scholar Top Publications resource<sup>4</sup>. To select journals that were explicitly 'continental'

or 'analytic', I relied upon a combination of self-identification (some journals explicitly identify themselves), the Leiter Report<sup>5</sup> (a listing of departments and journals principally voted upon by scholars in the field), and personal experience from presenting at continental philosophy conferences myself. In the second part, journals were taken exclusively from the top 10 of the Google Scholar Top Publications in the subject of philosophy, ostensibly to ensure that these are 'good' papers, which is short-hand for 'papers with a relatively large citation impact'. All 'bad' papers were taken from graduate and undergraduate philosophy journals, and it should be noted that authors of exceptional graduate papers are *actively* discouraged from publishing in such journals, as many scholars in the field believe publishing in such un-reputable format can have a negative impact on academic success<sup>6</sup>. The journals used were:

**Continental:** Philosophy and Phenomenological Research, European Journal of Philosophy, Continental Philosophy Review.

**Analytic:** Mind, Nous, The Philosophical Review, Analysis, The Journal of Philosophy.

**Good:** Mind, Philosophy and Phenomenological Research, The Philosophical Review, the APA journal, The Australasian Journal of Philosophy.

**Bad:** Episteme, Dianoia, Sapere Aude, Stance, Aporia, Ergo.

#### 3.2 Data Pre-processing

Since the bag-of-words approach must turn a text into numerical representations of term-frequency, a number of processing steps are required.

(i): The pdf files are read in and converted to strings using the python library *pdfminer*. While slow, this package is extremely good at extracting all the text from a .pdf file without errors. In every case the first page was omitted, since this is nearly always a title page including noisy data.

(ii): The data must be stemmed (which reduces similar words to a common stem) or lemmatized (which reduces words to their lemma based on a dictionary). Both were used, but lemmatization gave adequate results, and yielded more intelligible n-grams, since stemming tools often produce non-existent stems. Both processes were accomplished with the Natural Language Tool-kit for python, or *nlTK*.

(iii): All numbers and names were removed from the data, as this information would generally either be noise, or be a source of unwanted correlation. For example, in part 2, the journals often have information regarding dates in the footer or header of each page, whereas the student journals usually did not. Left in the data, this provided strong correlation with the mere presence of numbers, which showed up in the n-grams. Names were removed using *nlTK*, and numbers were removed using a substitution function (numbers were substituted with empty strings).

<sup>4</sup>"Google Scholar Top Publications." Google.com, scholar.google.com/citations?view\_op=top\_venues&hl=en&vq=hum.

<sup>5</sup>"Leiter Reports: A Philosophy Blog." Leiter Reports: A Philosophy Blog, leiterreports.typepad.com/.

<sup>6</sup>ibid.

(iv): Stop words, which are common words that don't convey any particular information, were removed using a built-in stop-word dictionary in the Sci-kit Learn python library, or *sklearn*. This dictionary had to be appended, however, with quite a few additional terms, again from the header and footer of the files. These terms created unwanted correlation, as mentioned above, since words like the journal names appear repeatedly in the training set. For this reason, an additional word list was created that removes all words from the headers and footers of each documents. This list includes things like journal names, URL prefixes, and months.

(v): Next, the cleaned text is divided into a training set, and a testing set for cross-validation, with a ratio of 85/15 for train/test.

vi: TF-IDF vectorization is then done using *sklearn*. The number of features used was varied to evaluate effectiveness. The options following options were tuned for accuracy, which will be discussed in the conclusion: max features, n-gram range, and max-df.

### 3.3 Classification

Compared to the pre-processing, classification was straightforward. The following classifiers are created using the *sklearn* package, trained, and then used on the test sets for cross-validation: K-nearest neighbors, Linear Discriminant Analysis, Quadratic Discriminant Analysis, and Support Vector Classification. Default options were used, and the  $n$  value in KNN was tuned for accuracy.

### 3.4 Visualization

Finally, the training and test sets were concatenated and decomposed using SVD to plot the singular values. For both data sets, a rank-2 representation of the data was plotted using the labels to give a sense of the distribution of the data.

## 4. Computational Results

The results from the experiment were excellent for KNN, with generally good results among the other methods. These values are collected in Figure 2. Singular values of the vectorized data show dominance of the first value, with a considerably long tail, shown in Figures 3 and 4. Scatter plots of the first two principal components show good distinction between the labeled classes, even with a rank-2 representation of the data, shown in Figures 5 and 6. Finally Table 2 shows the most correlated n-grams. This data was frustrating as I was not able to get the *chi2()* function to properly extract the class-specific n-grams, and so it must be noted that these n-grams are drawn from the entire corpus. Nevertheless, they were an important tool for identifying noise in the data.

## 5. Summary and Conclusions

The results suggest that we can classify philosophy texts based on high-level labels, in some cases with exceptionally high accuracy, both with KNN and SVC(in which the

	KNN	LDA	QDA	SVC
Part 1	100 / 96.7	80.8 / 96.7	100 / 50	100 / 100
Part 2	82.9 / 93.3	79.4 / 86.7	100 / 46.7	100 / 86.7

Figure 2. Classification results. Entries correspond to train/test accuracy.

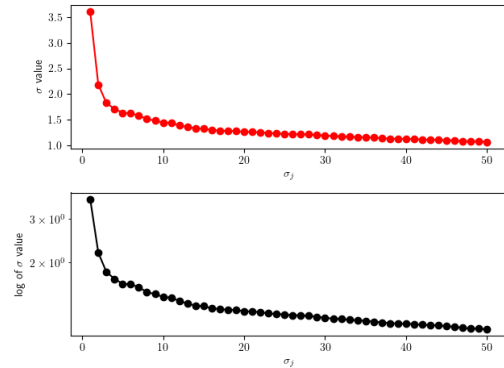


Figure 3. Continental v. Analytic, first 50 singular values.

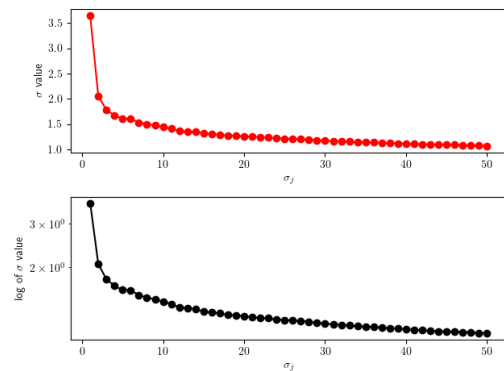


Figure 4. 'Good' v. 'Bad', first 50 singular values.

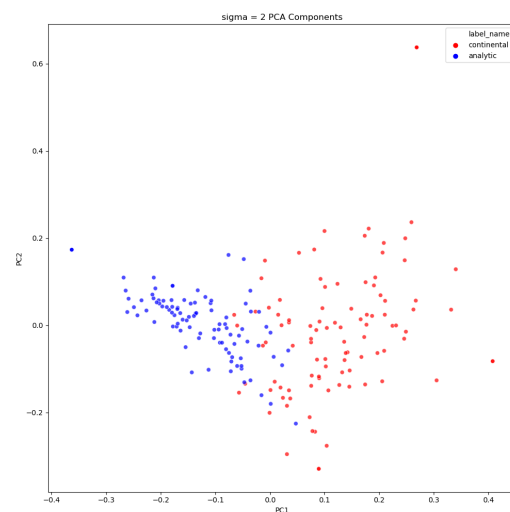


Figure 5. Part 1  $\sigma = 2$  Principal Components.

continental/analytic classification was perfect). This level of

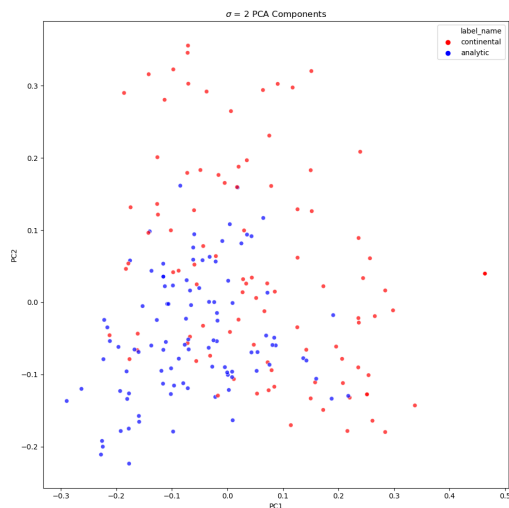


Figure 6. Part 1  $\sigma = 2$  Principal Components.

	Part 1	Part 2
Unigrams	operative	luck
	first	compatibilists
	fission	different
	tence	compatibilist
		determinism
Bigrams	true context	acprof oso
	palm tree	subject object
	account moral	theory The
	term body	epistemic reason
	doxastic obligation	classical logic

Table 2. Most correlated n-grams.

accuracy requires expert knowledge in a human, and so this result was quite surprising.

Parameter tuning had a significant impact on the accuracy of the results. Specifically, choosing a large number of features, 2000, showed Significant increases in accuracy over an initial trial of 500. Another impactful parameter was max-df, which discards any feature that appears in more than  $x$  documents. An initial value of 10 provided good results, but increasing it to 25 improved them substantially, likely because it included more unusual words that correlated strongly with the types of language used by the labeled groups. Finally, in the KNN classification, values of  $n$  were varied from 1 to 11. In part 1, a single neighbor provided optimal results, whereas in part 2 a large number of neighbors, 11, worked best.

The biggest issue with the experiment was the isolation of correlated n-grams. As someone with a significant philosophy background, this represented one of the most potentially interesting parts of the entire experiment, as it would have provided some sort of answer to the question: what type of language do the most impactful philosophers use? Future experimentation will focus on clearing up this bug in the code and trying to give concrete n-gram information.



## Appendix 1

`np.linalg.svd(A)`: Performs a singular value decomposition on the argument matrix.

`pdfminer.converter.TextConverter(rsrmgr, retstr)`: Creates an object to convert a pdf into a string. The standard resource manager is declared with `PDFResourceManager()`, and the standard string processor is declared with `io.StringIO()`. Full documentation for `pdfminer.six` can be found at: <https://pdfminersix.readthedocs.io/en/latest/>.

`os.walk(dir)`: Returns a string containing the root directory, a list containing the sub-directories as strings, and a list containing the files in the subdirectories as strings. Extremely useful for iterating over folder contents.

`nltk.SnowballStemmer("english")` and `WordNetLemmatizer()`: Both objects are used analogously. Declares a stemming/lemmatizing object using the natural language toolkit. A single word can be passed as a string using `stemmer.stem()`, and the object returns the stemmed or lemmatized word.

`nltk.word_tokenize(string)`: Tokenizes the passed string into a list of the individual words contained in the string. Important for stemming and lemmatization.

`nltk.tag.pos_tag()`: A function from the `nltk` that tags words based on types such as 'proper nouns'. Useful for processing strings based on word type. In particular it was used to remove all words with the tags for proper nouns and names.

`pd.DataFrame(dictionary)`: One of many methods for creating a pandas data frame object. In this case, a dictionary is passed and the data frame columns correspond to the dictionary values, with the keys serving as column titles.

`TfidfVectorizer()`: Creates an object to vectorize lists of strings and return inverse term-frequency data. A full discussion of the options can be found at [https://scikit-learn.org/stable/modules/generated/sklearn.feature\\_extraction.text.TfidfVectorizer.html](https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.TfidfVectorizer.html).

`KNeighborsClassifier(n_neighbors = n)`: Creates an object to classify data using the k-nearest neighbors method. Data is fitted and predicted using `classifier.fit()` and `classifier.predict()`. All the classifiers used have analogous syntax.

`(accuracy_score(labels, prediction))`: Returns the accuracy as a float from the prediction of an already trained classifier object. The 'prediction' parameter is the output of the `classifier.predict()` method.

`PCA(n_components = n)`: Returns an object that performs a singular value decomposition and returns the *n* principal components as a numpy array. The class method is used as follows: `PCA(n_components = n).fit_transform(matrix)`.

## Appendix 2

```

1 # Journal list:
2 # continental: philosophy and phenomenological research, european journal of philosophy, continental
   philosophy review
3 # analytic: mind, nous, the philosophical review, analysis, the journal of philosophy,
4 # good: mind, philosophy and phenomenological research, the philosophical review, the APA journal, the
   australasian journal of phil,
5 # bad: Episteme, dianoia, sapere aude, stance, aporia, ergo
6
7 ### Imports
8 import io
9 from pdfminer.pdfinterp import PDFResourceManager, PDFPageInterpreter
10 from pdfminer.converter import TextConverter
11 from pdfminer.layout import LAParams
12 from pdfminer.pdfpage import PDFPage
13 import os
14 import pickle
15 import pandas as pd
16 import re
17 import nltk
18 from nltk.corpus import stopwords
19 from nltk.stem import WordNetLemmatizer

```

```

20 from sklearn.feature_extraction.text import TfidfVectorizer
21 from sklearn.model_selection import train_test_split
22 from sklearn.feature_selection import chi2
23 import numpy as np
24 import matplotlib.pyplot as plt
25 import seaborn as sns
26
27 ### Functions for the project
28
29 # Perform a reduced SVD of the data for Part 1 and plot the singular values on a standard and semi-log
    axis.
30 def svd_plot(data):
31     A1 = data
32
33     U, S, V = np.linalg.svd(A1, full_matrices=False)
34     x = np.linspace(1, 50, 50)
35
36     # Plots the first 50 singular values.
37     fig, (ax1, ax2) = plt.subplots(2)
38     ax1.plot(x, S[0:50], 'r-o')
39     ax1.set_xlabel('$\sigma_j$')
40     ax1.set_ylabel('$\sigma$ value')
41
42     ax2.semilogy(x, S[0:50], 'k-o', )
43     plt.rc('text', usetex=True)
44     ax2.set_xlabel('$\sigma_j$')
45     ax2.set_ylabel('log of $\sigma$ value')
46     plt.show()
47
48 # Function that takes the path of a pdf and returns it as a string object.
49 def convert_pdf_to_txt(path):
50     rsrcmgr = PDFResourceManager()
51     retstr = io.StringIO()
52     codec = 'utf-8'
53     laparams = LAParams()
54     device = TextConverter(rsrcmgr, retstr, codec=codec, laparams=laparams)
55     fp = open(path, 'rb')
56     interpreter = PDFPageInterpreter(rsrcmgr, device)
57     password = ""
58     maxpages = 0
59     caching = True
60     pagenos=set()
61
62     # Skips the first page of the pdf file since it contains data that can be considered noise.
63     pages = PDFPage.get_pages(fp, pagenos, maxpages=maxpages, password=password, caching=caching,
        check_extractable=True)
64     iter_pages = iter(pages)
65     next(iter_pages)
66     for page in iter_pages:
67         test = page
68         interpreter.process_page(page)
69
70     text = retstr.getvalue()
71     fp.close()
72     device.close()
73     retstr.close()
74     return text
75
76 # A function that iterates through a given root directory and extracts all pdf files as a list of
    strings.
77 def pdf_to_strings(path):
78     rootdir = path
79     corpus_raw = []
80     index = 0
81     label_list = []
82
83     # Iterates through the base folder with os.walk, which goes through each subfolder, with
84     # iterate variables for the files, the subdirectories and the directories.
85     for subdir, dirs, files in os.walk(rootdir):
86

```

```

87     # Iterates through each file.
88     for file in files:
89
90         # Gets the file path to pass to the extraction function.
91         file_path = os.path.join(subdir, file)
92         # Appends the name of the current folder, which is the label for the pdf
93         label_list.append(os.path.basename(os.path.normpath(subdir)))
94         corpus_raw.append(convert_pdf_to_txt(file_path))
95     return corpus_raw, label_list;
96
97 # Function to generate all the raw lists of strings for each pdf. Significant Processing time.
98 def process_paths(path1, path2):
99     corpus1, label_list1 = pdf_to_strings(path1)
100    corpus2, label_list2 = pdf_to_strings(path2)
101    return corpus1, label_list1, corpus2, label_list2
102
103 # A function to stem a list of strings composed of pdfs. Does so iteratively, then joins them back
104 # into single strings and returns a list of strings like the argument.
105 def list_stemmer(doc_list):
106     n = len(doc_list)
107     stemmed_text_list = []
108
109     # Initialize the stemmer
110     stemmer = nltk.SnowballStemmer("english")
111
112     for i in range(0, n):
113         # According to stackexchange discussions, a list comprehension is much faster for this task
114         # than a loop.
115         stemmed_text = ' '.join(stemmer.stem(token) for token in nltk.word_tokenize(doc_list[i]))
116         stemmed_text_list.append(stemmed_text)
117     return stemmed_text_list
118
119 def list_lemmatizer(doc_list):
120     from nltk.stem import WordNetLemmatizer
121     n = len(doc_list)
122     lemmad_text_list = []
123
124     # Initialize the lemmatizer
125     lemmatizer = WordNetLemmatizer()
126
127     for i in range(0, n):
128         # According to stackexchange discussions, a list comprehension is much faster for this task
129         # than a loop.
130         lemmad_text = ' '.join(lemmatizer.lemmatize(token) for token in nltk.word_tokenize(doc_list[i]))
131         lemmad_text_list.append(lemmad_text)
132     return lemmad_text_list
133
134 # Removes numbers from the argument, which should be a list of strings.
135 def num_removal(doc_list):
136     for i, list in enumerate(doc_list):
137         doc_list[i] = re.sub(r'\d+', '', list)
138
139 # Removes most names from the argument, which should be a list of strings
140 def name_removal(doc_list):
141     for i, list in enumerate(doc_list):
142         tagged_string = nltk.tag.pos_tag(list.split())
143         new_string = [word for word, tag in tagged_string if tag != 'NNP' and tag != 'NNPS']
144         doc_list[i] = ' '.join(new_string)
145
146 # Load the desired data. Arguments are the path and whether raw, stemmed, or lemmatized data
147 # is desired. Returns part1_data, part1_labels, part2_data, part2_labels. Code has no error
148 # control, type should be "lemmed", "stemmed" or "raw"
149 def load_data(path, type):
150     import pickle
151     with open(path + "/part1_data_" + type, 'rb') as f:
152         part1_data = pickle.load(f)
153
154     with open(path + "/part1_labels", 'rb') as f:
155         part1_labels = pickle.load(f)
156

```



```

157     with open(path + "/part2_data_" + type, 'rb') as f:
158         part2_data = pickle.load(f)
159
160     with open(path + "/part2_labels", 'rb') as f:
161         part2_labels = pickle.load(f)
162     return part1_data, part1_labels, part2_data, part2_labels;
163
164 # Save files based on type. Again, no error control, same type as the load function
165 # must be used.
166 def save_data(path, type, corpus1, label_list1, corpus2, label_list2):
167     with open(path + "/part1_data_" + type, 'rb') as f:
168         pickle.dump(corpus1, f)
169
170     with open(path + "/part1_labels", 'rb') as f:
171         pickle.dump(label_list1, f)
172
173     with open(path + "/part2_data_" + type, 'rb') as f:
174         pickle.dump(corpus2, f)
175
176     with open(path + "/part2_labels", 'rb') as f:
177         pickle.dump(label_list2, f)
178
179 ### Download nltk resources and load the various functions written for the project from file. Only
180 really
181 # needs to be run once.
182 nltk.download('stopwords')
183 nltk.download('punkt')
184 nltk.download('wordnet')
185 nltk.download('averaged_perceptron_tagger')
186
187 ### Get the raw strings and pickle the data
188 path1 = "C:/Users/zennsunni/Dropbox/School Stuff/Winter 2020/AMATH.582/Project/part1"
189 path2 = "C:/Users/zennsunni/Dropbox/School Stuff/Winter 2020/AMATH.582/Project/part2"
190 corpus1, label_list1, corpus2, label_list2 = process_paths(path1, path2)
191
192 path = "C:/Users/zennsunni/Dropbox/School Stuff/Winter 2020/AMATH.582/Project"
193 type = "raw"
194 save_data(path, type, corpus1, label_list1, corpus2, label_list2)
195
196 ### Stem the data and save these files separately
197 corpus1 = list_stemmer(part1_data_raw)
198 corpus2 = list_stemmer(part2_data_raw)
199
200 save_data(path, type, corpus1, label_list1, corpus2, label_list2)
201
202 ### Lemmatize the data and save these files separately
203 corpus1 = list_lemmatizer(part1_data_raw)
204 corpus2 = list_lemmatizer(part2_data_raw)
205
206 save_data(path, type, corpus1, label_list1, corpus2, label_list2)
207
208 ### Load the desired data
209 path = "C:/Users/zennsunni/Dropbox/School Stuff/Winter 2020/AMATH.582/Project"
210 type = "lemmed"
211 part1_data, part1_labels, part2_data, part2_labels = load_data(path, type)
212
213 ### Remove numbers and most names from data before setup
214 num_removal(part1_data)
215 num_removal(part2_data)
216
217 name_removal(part1_data)
218 name_removal(part2_data)
219
220 ### Setup and Classification
221 from sklearn.feature_extraction import text
222
223 data1 = part1_data
224 data2 = part2_data
225

```

```

226 # This dictionary gives the folder names as keys to the artists inside them.
227 label_index1 = {"continental": 0, "analytic":1}
228 label_index2 = {"good": 0, "bad": 1}
229
230 # Create numerical labels
231 part1_labels_int = [label_index1[q] for q in part1_labels]
232 part2_labels_int = [label_index2[q] for q in part2_labels]
233
234 # Put the data in a data frame
235 prep_dict1 = {'part1_data': data1, 'part1_labels':part1_labels, 'part1_id': part1_labels_int}
236 prep_dict2 = {'part2_data': data2, 'part2_labels':part2_labels, 'part2_id': part2_labels_int, }
237 df1 = pd.DataFrame(prepare_dict1)
238 df2 = pd.DataFrame(prepare_dict2)
239
240 # Split the corpus up into randomized training/testing sets
241 X_train1, X_test1, y_train1, y_test1 = train_test_split(df1['part1_data'], df1['part1_id'], test_size
    =0.15, random_state=8)
242 X_train2, X_test2, y_train2, y_test2 = train_test_split(df2['part2_data'], df2['part2_id'], test_size
    =0.15, random_state=8)
243
244 #%%
245
246 # Parameter selection for vectorization of the data. We will look at unigrams, bigrams, and trigrams. We
    will
247 # Also ignore words that appear in less than 5% of the documents. Max features is set very high, since a
    few odd words
248 # Might be very correlated with a particular label, and philosophy papers tend to use the full spectrum
    of the
249 # English Language. This might have to be tuned up to a very high number.
250 ngram_range = (1,2)
251 min_df = 1
252 max_df = 25
253 max_features = 2000
254
255 # Additional stop words, which get added to sklearn's stop word list.
256 stop_list = ["Authors", "Published", "Permissions", "journal", "journals", "reserved", "Inc", "
    Philosophical", "Review",
257             "org", "This", "content", "downloaded", "https", "All", "Oxford", "New", "York", "
    university", "University", "Vol", "http", "doi", "Press",
258             "And", "org", "UTC", "Mar", "ed", "downloaded", "Sun", "Mon", "Tues", "Wed", "Thur", "Fri",
    "jstor", "Cornell",
259             "DOI", "European", "Continental", "Springer", "Sons", "Wiley", "wileyonlinelibrary",
    "January", "February",
260             "March", "April", "May", "June", "July", "August", "September", "October", "November", "
    December", "Stance",
261             "Dianoia", "Undergraduate", "Boston", "College", "Aporia", "Ergo", "no." "no", "
    Phenomenological", "Research",
262             "LLC", "Australasian", "Routledge", "Association", "pp", "In", "html"]
263 stop_words = text.ENGLISH_STOP_WORDS.union(stop_list)
264
265 # Declare the vectorizer. Note this is a TF-IDF vectorizer, and thus weights low-frequency words,
266 # which is perfect for analyzing the very tribal lexicons used by philosophers.
267 tfidf = TfidfVectorizer(encoding='utf-8',
268                         ngram_range=ngram_range,
269                         stop_words= stop_words,
270                         lowercase=False,
271                         max_df=max_df,
272                         min_df=min_df,
273                         max_features=max_features,
274                         norm='l2',
275                         sublinear_tf=True)
276
277 # Extract the features for the training and test sets for Part 1.
278 features_train1 = tfidf.fit_transform(X_train1).toarray()
279 labels_train1 = y_train1
280 features_test1 = tfidf.transform(X_test1).toarray()
281 labels_test1 = y_test1
282
283 # Extract the features for the training and test sets for Part 2.
284 features_train2 = tfidf.fit_transform(X_train2).toarray()

```

```

285 labels_train2 = y_train2
286 features_test2 = tfidf.transform(X_test2).toarray()
287 labels_test2 = y_test2
288
289 ### Plot the singular values
290
291 features_full1 = np.concatenate((features_train1, features_test1), axis=0)
292 features_full2 = np.concatenate((features_train2, features_test2), axis=0)
293 ###
294 svd_plot(features_full2)
295
296
297 ### Prints the most correlated uni/bi/trigrams
298
299 from sklearn.feature_selection import chi2
300
301 for field, index in sorted(label_index1.items()):
302     features_chi2 = chi2(features_train1, labels_train1 == index)
303     indices = np.argsort(features_chi2[0])
304     feature_names = np.array(tfidf.get_feature_names())[indices]
305     unigrams = [v for v in feature_names if len(v.split(' ')) == 1]
306     bigrams = [v for v in feature_names if len(v.split(' ')) == 2]
307     # trigrams = [v for v in feature_names if len(v.split(' ')) == 3]
308     print("# '{}' category:".format(field))
309     print(" . ### Most correlated unigrams:\n. {}".format('\n. '.join(unigrams[-5:])))
310     print(" . ### Most correlated bigrams:\n. {}".format('\n. '.join(bigrams[-5:])))
311     # print(" . ### Most correlated trigrams:\n. {}".format('\n. '.join(trigrams[-5:])))
312     print("")
313
314 ### Classify with KNN, Part 1
315 from sklearn.neighbors import KNeighborsClassifier
316 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
317
318 neigh = KNeighborsClassifier(n_neighbors=1)
319 neigh.fit(features_train1, labels_train1)
320 predict1 = neigh.predict(features_test1)
321
322 # Training accuracy
323 print("The training accuracy is: ")
324 print(accuracy_score(labels_train1, neigh.predict(features_train1)))
325
326 # Test accuracy
327 print("The test accuracy is: ")
328 print(accuracy_score(labels_test1, predict1))
329
330 ### Classify with KNN, Part 2
331 from sklearn.neighbors import KNeighborsClassifier
332 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
333
334 neigh = KNeighborsClassifier(n_neighbors=9)
335 neigh.fit(features_train2, labels_train2)
336 predict2 = neigh.predict(features_test2)
337
338 # Training accuracy
339 print("The training accuracy is: ")
340 print(accuracy_score(labels_train2, neigh.predict(features_train2)))
341
342 # Test accuracy
343 print("The test accuracy is: ")
344 print(accuracy_score(labels_test2, predict2))
345
346 ### SVC Fit for comparison, Part1
347 from sklearn import svm
348 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
349
350 svm_clf = svm.SVC()
351 svm_clf.fit(features_train1, labels_train1)
352 predict1 = svm_clf.predict(features_test1)
353
354 # Training accuracy

```

```

355 print("The training accuracy is: ")
356 print(accuracy_score(labels_train1 , svm_clf.predict(features_train1)))
357
358 # Test accuracy
359 print("The test accuracy is: ")
360 print(accuracy_score(labels_test1 , predict1))
361
362 ### SVC Fit for comparison , Part2
363 from sklearn import svm
364 from sklearn.metrics import classification_report , confusion_matrix , accuracy_score
365
366 svm_clf = svm.SVC()
367 svm_clf.fit(features_train2 , labels_train2)
368 predict2 = svm_clf.predict(features_test2)
369
370 # Training accuracy
371 print("The training accuracy is: ")
372 print(accuracy_score(labels_train2 , svm_clf.predict(features_train2)))
373
374 # Test accuracy
375 print("The test accuracy is: ")
376 print(accuracy_score(labels_test2 , predict2))
377
378 ### LDA for comparison Part 1
379 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
380
381
382 lda = LDA(n_components=1)
383 lda.fit(features_train1 , labels_train1)
384 predict1 = lda.predict(features_test1)
385
386 # Training accuracy
387 print("The training accuracy is: ")
388 print(accuracy_score(labels_train1 , lda.predict(features_train1)))
389
390 # Test accuracy
391 print("The test accuracy is: ")
392 print(accuracy_score(labels_test1 , predict1))
393
394 ### LDA for comparison Part 2
395
396 lda = LDA(n_components=1)
397 lda.fit(features_train2 , labels_train2)
398 predict2 = lda.predict(features_test2)
399
400 # Training accuracy
401 print("The training accuracy is: ")
402 print(accuracy_score(labels_train2 , lda.predict(features_train2)))
403
404 # Test accuracy
405 print("The test accuracy is: ")
406 print(accuracy_score(labels_test2 , predict2))
407
408 ### QDA for comparison Part 1
409 from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis as QDA
410
411 qda = QDA()
412 qda.fit(features_train1 , labels_train1)
413 predict1 = qda.predict(features_test1)
414
415 # Training accuracy
416 print("The training accuracy is: ")
417 print(accuracy_score(labels_train1 , qda.predict(features_train1)))
418
419 # Test accuracy
420 print("The test accuracy is: ")
421 print(accuracy_score(labels_test1 , predict1))
422
423 ### QDA for comparison Part 2
424

```

```

425
426 qda = QDA()
427 qda.fit(features_train2 , labels_train2)
428 predict2 = qda.predict(features_test2)
429
430 # Training accuracy
431 print("The training accuracy is: ")
432 print(accuracy_score(labels_train2 , qda.predict(features_train2)))
433
434 # Test accuracy
435 print("The test accuracy is: ")
436 print(accuracy_score(labels_test2 , predict2))
437
438
439
440 ### Visualize the features in a plot
441 from sklearn.decomposition import PCA
442 import matplotlib.pyplot as plt
443 import seaborn as sns
444
445 features = np.concatenate((features_train1 , features_test1) , axis=0)
446 labels = np.concatenate((labels_train1 , labels_test1) , axis=0)
447 title = "sigma = 2 PCA Components"
448 princ_comps = PCA(n_components=2).fit_transform(features)
449
450 # Put them into a dataframe
451 df_features = pd.DataFrame(data= princ_comps ,
452                             columns=['PC1' , 'PC2'])
453
454 # Now we have to paste each row's label and its meaning
455 # Convert labels array to df
456 df_labels = pd.DataFrame(data=labels ,
457                             columns=['label'])
458
459 df_full = pd.concat([df_features , df_labels] , axis=1)
460 df_full['label'] = df_full['label'].astype(str)
461
462 # Makes a new dictionary that is flipped , to unzip the label codes the other
463 # direction .
464 new_labels = {"0": "continental" , "1": "analytic"}
465
466 # And map labels
467 df_full['label_name'] = df_full['label']
468 df_full = df_full.replace({'label_name': new_labels})
469
470 plt.figure(figsize=(10 , 10))
471 sns.scatterplot(x='PC1' ,
472                 y='PC2' ,
473                 hue='label_name' ,
474                 data=df_full ,
475                 palette=["red" , "blue"] ,
476                 alpha=.7).set_title(title);
477
478 plt.savefig('part1_scatter.png' , facecolor = "white")
479 plt.show()
480
481 ### And for part 2
482
483 features = np.concatenate((features_train2 , features_test2) , axis=0)
484 labels = np.concatenate((labels_train2 , labels_test2) , axis=0)
485 title = "\sigma = 2 PCA Components"
486 princ_comps = PCA(n_components=2).fit_transform(features)
487
488 # Put them into a dataframe
489 df_features = pd.DataFrame(data= princ_comps ,
490                             columns=['PC1' , 'PC2'])
491
492 # Now we have to paste each row's label and its meaning
493 # Convert labels array to df
494 df_labels = pd.DataFrame(data=labels ,

```

```

495         columns=['label'])
496
497 df_full = pd.concat([df_features, df_labels], axis=1)
498 df_full['label'] = df_full['label'].astype(str)
499
500 # Makes a new dictionary that is flipped, to unzip the label codes the other
501 # direction.
502 new_labels = {"0": "continental", "1": "analytic"}
503
504 # And map labels
505 df_full['label_name'] = df_full['label']
506 df_full = df_full.replace({'label_name': new_labels})
507
508 plt.figure(figsize=(10, 10))
509 sns.scatterplot(x='PC1',
510                y='PC2',
511                hue="label_name",
512                data=df_full,
513                palette=["red", "blue"],
514                alpha=.7).set_title('title');
515
516 plt.savefig('part2-scatter.png', facecolor = "white")
517 plt.show()

```